

CS61C Summer 2012 Midterm

Your Name: _____ SID: _____

Your TA (Circle): Raphael Paul Brandon Sung Roa

Name of person to your LEFT: _____

Name of person to your RIGHT: _____

This exam is worth 110 points and will count for 20% of your course grade.

The exam contains 9 questions on 13 numbered pages, including the cover page. Put all answers on these pages; don't hand in stray pieces of paper.

Question 0: You will receive 1 point for properly filling out this page as well your login on every page of the exam.

Question	Points (Minutes)	Score
0	1 (0)	
1	15 (26)	
2	10 (14)	
3	10 (18)	
4	9 (16)	
5	10 (18)	
6	17 (26)	
7	12 (22)	
8	26 (40)	
Total	110 (180)	

All the work is my own. I had no prior knowledge of the exam contents nor will I share the contents with others in CS61C who have not taken it yet.

Signature: _____

Question 1: *Potpourri – Hard to spell, nice to smell...* (15 points, 26 minutes)

a) **True/False:**

- T F We prefer two's complement over the unsigned representation because two's complement can represent more values.
- T F The assembler uses symbol tables to resolve absolute addresses.
- T F A program will always execute faster in a RISC architecture than a CISC architecture.
- T F The greater the number of memory accesses in a program, the greater the AMAT.
- T F Pseudo-instructions do not always use `$at`.
- T F Since `$s0` is a "saved register," it does not need to be saved before any function calls.
-

b) Fill in the function below, which returns a new copy of the argument (struct definition not shown):

```
struct something *cpySomething(struct something *old) {
```

```
    _____  
    _____  
    return new;  
}
```

c) Here is what is currently on the heap:

A	A	A			C	D	D		
---	---	---	--	--	---	---	---	--	--

The order of allocation/frees: A allocated, B allocated, C allocated, B freed, D allocated

What allocation strategy was used? _____

d) What is the average CPI of the program described in the table to the right? Which is better: halving memory access cycles or arithmetic cycles and why?

Instruction Category	Cycles	Frequency
Memory Access	10	0.1
Arithmetic	2	0.4
Branch	3	0.2
Comparison	1	0.3

- e) I'm getting the message "cannot execute binary file". In one or two sentences, explain what the problem is and how to fix it.

- f) In our 32-bit single-precision floating point representation, we decide to convert one significant bit to an exponent bit. How many **denormalized numbers** do we have relative to before? (Circle one)

More

Fewer

Rounded to the nearest power of 2, how many denorm numbers are there in our new format?
(Answer in IEC format)

Question 2: *Flippin' Fo' Fun* (10 points, 14 minutes)

Assume that the most significant bit (MSB) of x is a 0. We store the result of flipping x 's bits into y . Interpreted in the following number representations, how large is the magnitude of y relative to the magnitude of x ? Circle ONE choice per row.

Unsigned	$ y < x $	$ y = x $	$ y > x $	Can't Tell
One's Complement	$ y < x $	$ y = x $	$ y > x $	Can't Tell
Two's Complement	$ y < x $	$ y = x $	$ y > x $	Can't Tell
Sign and Magnitude	$ y < x $	$ y = x $	$ y > x $	Can't Tell
Biased Notation (e.g. FP exponent)	$ y < x $	$ y = x $	$ y > x $	Can't Tell

Question 3: *Doctor Who?!?* (10 Points, 18 Minutes)

The Daleks are invading the Earth again, and we need the help of the Doctor! Find the errors in this code and fix them so that the code correctly prints **"The 10th Doctor and the Blue Police Box"**. There is exactly one coding error for each function and function call pair and can be fixed by changing 5 or 6 lines total. Fill in the corrections in the blanks on the opposite page.

```
1    void whichDoctor(int* input) {
2        input = 10;
3    }

4    void doctorChanger(char** input1, char** input2) {
5        char* temp = *input1;
6        *input1 = *input2;
7        *input2 = temp;
8    }

9    char* policeBoxGiver(char* input) {
10        *input = "The Master";
11        return "Police Box";
12    }

13    char* colorMaker(void) {
14        char* color = malloc(sizeof(char) * 4);
15        color[0] = 'B';
16        color[1] = 'l';
17        color[3] = 'u';
18        color[2] = 'e';
19        color[4] = 0;
20        return color;
21    }

22    char* colorFixer(char* input) {
23        char temp = *(input+2);
24        *(input+2) = *(input+1);
25        *(input+1) = temp;
26    }

27    int main(void) {
28        int * ith = malloc(sizeof(int));
29        whichDoctor(ith);
30        char* doctor = "Master";
31        char* master = "Doctor";
32        char* details = "and the";
33        char* color = colorMaker();
34        colorFixer(color);
35        char* box = "David Tennant";
36        doctorChanger(doctor, master);
37        policeBoxGiver(box);
38        printf("The %dth %s %s %s %s", *ith, doctor, details, color, box);
39    }
```

Line #	Corrected Code
_____	_____
_____	_____
_____	_____
_____	_____
_____	_____
_____	_____

Question 4: *Let Me Float This Idea By You* (9 Points, 16 Minutes)

For a very simple household appliance like a thermostat, a more minimalistic microprocessor is desired to reduce power consumption and hardware costs. We have selected a **16-bit** microprocessor that does not have a floating-point unit, so there is no native support for floating point operations (no `float/double`). However, we'd still like to represent decimals for our temperature reading so we're going to implement floating point operations in software (in C).

a) Define a new variable type called `fp`:

We have decided to use a representation with a **5-bit exponent field** while following all of the representation conventions from the MIPS 32-bit floating point numbers **except denorms**.

Fill in the following functions. Not all blanks need to be used. You can call these functions and assume proper behavior regardless of your implementation. Assume our hardware implements the C operator ">>" as *shift right arithmetic*.

b)

```
/* returns -num */
fp negateFP(fp num) {
    return _____;
}
```

c)

```
/* returns the signed value of the exponent */
int getExp(fp num) {
    _____
    return _____;
}
```

d)

```
/* multiplies floating point num by 2^n, while detecting over/underflow */
/* remember, there are no denorms */
fp multPow2(fp num, int n) {
    _____
    if(_____) exit(1); #overflow
    if(_____) exit(-1); #underflow
    _____
    return _____;
}
```

Question 5: *Who Says Less is Better?* (10 points, 18 Minutes)

We're going to take a page out of the ARM book and design a new instruction set architecture with just **16 32-bit registers**. This means that we only need 4-bit register fields in our instructions.

- a) How many extra bits do we have now for other fields in the following formats?

R: _____ J: _____

- b) For R-format instructions, would you give the extra bits to `opcode`, `shamt`, or `funct`? _____
Explain your choice in a sentence or two (no credit without explanation):

For I-format instructions, we naturally give the extra bits to the `immediate` field, resulting in the following format:

[`opcode` (6) | `rs` (4) | `rt` (4) | `immediate` (18)]

- c) What fraction of our address space can we now reach with a branch instruction?

- d) Assume our PC currently contains the address `0x08000000`.

What is the LOWEST address (in hex) we can reach with a branch?

- e) Write out the Verilog pseudocode (as in the OPERATION column on the MIPS Green Sheet) for `beq`. Make sure you specify what `BranchAddr` is.

Question 6: *Cache in While You Can* (17 points, 26 Minutes)

Consider a single 4KiB cache with 512B blocks and a write-back policy. Assume a 32-bit address space.

a) If the cache were direct-mapped,

of rows? _____ # of offset bits? _____

b) If the cache were 4-way set associative,

of tag bits? _____ # of index bits? _____ # of bits per cache slot? _____

Consider an array of the following `location` structs:

```
typedef struct {  
    ... // some undefined number of other struct members  
    int visited;  
    int danger;  
} location;  
  
location locs[NUM_LOCS];
```

Here's a piece of code that counts the number of places we've visited. Assume this gets executed somewhere in the middle of our program, that `count` is held in a register, and the size of the array is greater than 4 KiB.

```
for(int i = 0; i < NUM_LOCS; i++)  
    if(locs[i].visited) count++;
```

c) What's the fewest possible number of bytes written to main memory? _____

d) What's the greatest possible number of bytes written to main memory? _____

Now consider if we store the `visited` and `danger` information in individual arrays instead:

```
int visited[NUM_LOCS];  
int danger[NUM_LOCS];
```

e) This way, the cache can exploit better _____ for the above task.

We can expect a _____ (higher or lower) miss rate

because of the change in the number of _____ (type of cache miss) misses.

Consider the following code with `NUM_LOCS > 210`.

```
for(int i = 0; i < NUM_LOCS; i++)  
    if(visited[i] && danger[i] > 5) count++;
```

Two memory accesses are made per iteration: one into `visited`, the other into `danger`. Assume that the cache has no valid blocks initially. **You are told that in the worst case, the cache has a miss rate of 100%.** Consider each of the following possible changes to the cache individually.

f) Mark each as **E**, if it eliminates the chances of this worst-case scenario miss rate, **R** if it reduces the chances, or **N** if it's not helpful.

- More sets, same block size, same associativity _____
- Double associativity, half block size, same total cache size _____
- Everything stays the same but use a write-through policy instead _____

Question 7: *Can't Make Copies Fast Enough* (12 points, 22 Minutes)

We are revisiting our friend the Fast String Copy from lecture! Recall that the function prototype in C is as follows:

```
char *strcpy(char *dst, char *src);
```

Consider the following MIPS implementation of this function:

```
jal  strcpy  # begin function call
...
strcpy:
    addi  $v0,$a0,0
loop: lb   $t0,0($a1)
    sb   $t0,0($a0)
    addiu $a0,$a0,1
    addiu $a1,$a1,1
    beq  $t0,$zero,exit
    j    loop
exit: jr   $ra
```

Suppose we are running code on a machine with the following cache parameters:

- **Unified** L1\$ with a hit time of 2 cycles and a hit rate of 95%
- Miss Penalty to main memory of 200 cycles
- Base CPI of 1.5 (in the absence of cache misses)

a) Calculate our machine's AMAT:

b) What is the CPI of a single call to `strcpy` with `src = ""` (the function call includes the `jal`)?

c) We decide to add a L2\$ to reduce our AMAT to 6. Our L2\$ has a hit time of 20 cycles. What's the worst Local Hit Rate that will still meet our AMAT goal?

- d) In addition to speeding up our architecture, we want to speed up our code, so we decide to eliminate the return value (presumably the caller retains a copy of the destination pointer). In this case, the `strcpy` function above can be rewritten in just 6 instructions. Write out this implementation in the blanks below, introducing any necessary labels (don't worry about any label name clashes with `strcpy`).

`strcpy2:`

- e) If we call `strcpy` and `strcpy2` on the same `src` string of length $n+1=N$ (N includes `'\0'`), what is the ratio of instructions executed in `strcpy` to instructions executed in `strcpy2` (including the `jal`)? Leave your answer in terms of N .
- f) Is the ratio in part (e) the same as the relative performance between these two functions? In a sentence or two, explain why or why not.


Question 8: *Putting the Science in Computer Science* (26 points, 40 minutes)

DNA can be called the “alphabet of life.” From a very simplified view, DNA within a cell produces amino acids, which in turn produce proteins, which are the building blocks for most of your body. Here we’d like to write some code for examining a strand of DNA.

- a) DNA is made up of *nucleotides*, which we write shorthand as A, C, G, and T. DNA is in base 4 (quaternary)! Fill in the table below, using the DNA nucleotide symbols in alphabetical order ($A < C < G < T$).

Decimal	DNA
	CAT
50	

An amino acid is encoded by three nucleotides. Because DNA is found in long strands, the following 5 nucleotides can be read 3 different ways:



The sequence **ATG** (as seen in the 3rd reading frame) signals the beginning of a protein (“start codon”).

- b) Fill in the blanks on the opposite page for the **recursive** function `find_start` in MIPS that returns the position of the first start codon found in the given strand of DNA. Assume each nucleotide is stored as a `char` in memory. *Blanks do not necessarily need to be filled*. Maximum points awarded for using the *fewest* amount of registers and memory.

[Answer the following AFTER looking at the code]

Assume we call `find_start` from main with `char dna[] = "GCATGC";`.

- c) How many total frames are created on the Stack (not including `main`)? _____
- d) What is the maximum depth of the Stack (in # of frames, not including `main`)? _____
- e) What will the line `j ret` look like once this file is run through the assembler? _____
- f) Where will the label `ret` show up? (Circle one)

Symbol Table

Relocation Table

Both

Neither

C function prototype: `/* dna: start address of DNA strand */`

`/* pos: search position from start of strand */`

`int find_start(char *dna, int pos);`

this is a recur. function don't forget dummy

```

find_start:
    addiu $sp,$sp,-4 # PROLOGUE
    sw $ra, 0($sp) #don't need to set up par. since dna stays the same for jal call

    jal     strlen     # call strlen(dna); Assume strlen doesn't
                        # change $a0 or $a1
    slti $t0, $v0, 3    # make sure we don't read past the end of
                        # the array

    beq     $t0,$0,chk  # 'chk' for check if start codon
    addi    $v0,$0,-1   # return -1 (start codon not found)
    j       ret         # 'ret' for return

chk:      Text
    lb      $t1,0($t0)
    addi    $t2,$0,65   why do you add 65?
    bne     $t1,$t2,rec # 'rec' for recurse
    lb      $t1,1($t0)
    addi    $t2,$0,84
    bne     $t1,$t2,rec
    lb      $t1,2($t0)
    addi    $t2,$0,71
    bne     $t1,$t2,rec
    add $v0, $t2, $0    # return current position
    j       ret

rec:      lb $t1, 4($a1) # recurse at next position
    addi $t1, $t1, 1
    Text
    jal  find_start
    Text
    Text

ret:      # EPILOGUE
    sw $ra, 0($sp)
    addiu $sp,$sp,4
    jr    $ra

```

BACK OF EXAM

(Any work on this page will not be graded)